MLOps Assignment 1 Report - Surbhi Arya

Heart Risk Prediction — End-to-End MLOps Pipeline 1
Overview 1
Part I — Problem & Data 1
1. Objective 2
2. Dataset & EDA 2
Part II — Modelling 4
3. Feature Engineering 5
4. Model Training & Results 5
Results 6
5. Experiment Tracking (MLflow) 9
6. Model Packaging & Reproducibility 11
Part III — Engineering & Operations 11
7. Testing Strategy 11
8. CI/CD Pipeline 12
9. Containerization 12
10. Production Deployment (Kubernetes via Helm) 13
11. Monitoring & Logging 14
Appendix 15
A. Setup Instructions 16
B. Architecture Diagram 16
C. Repository Link 19

Heart Risk Prediction — End-to-End MLOps Pipeline

Course	AIMLCZG523
Course Name	Machine Learning Operations (MLOps)
Assignment	01 - End-to-End ML Model Development, CI/CD, and Production Deployment
Author	Surbhi Arya
BITS ID	2025AE05688
Repository	https://github.com/arya2305/mlops-workflow
Video Recording Link	2025AE05688_mlops-workflow.mov

Overview

An end-to-end, reproducible pipeline for a binary classifier predicting heart disease risk from patient vitals — covering data ingestion, EDA, feature engineering, model training with experiment tracking, automated testing, CI/CD, containerization, Kubernetes deployment via Helm, and monitoring.

Every stage described in this report was actually executed during development: the dataset was pulled live from UCI, both models were trained and tracked in MLflow (with the winner registered in the Model Registry), the full `unittest` suite was run, the Docker image was built and smoke-tested, and the complete stack — API, Ingress, Prometheus, and Grafana — was deployed to a local Minikube cluster with a single `helm install` and verified against real traffic. Screenshots for each stage are in `screenshots/`.

This report is organized into three parts followed by an appendix:

- **Part I — Problem & Data:** the dataset and what the EDA revealed
- **Part II — Modelling:** feature engineering, model training/results, and experiment tracking
- **Part III — Engineering & Operations:** testing, CI/CD, containerization, deployment, and monitoring
- **Appendix:** setup instructions, architecture diagram, repository link

---

Part I — Problem & Data

1. Objective

Build a classifier that predicts heart disease risk from patient vitals, packaged and operated the way a small production service actually would be — not just a notebook that produces a number.

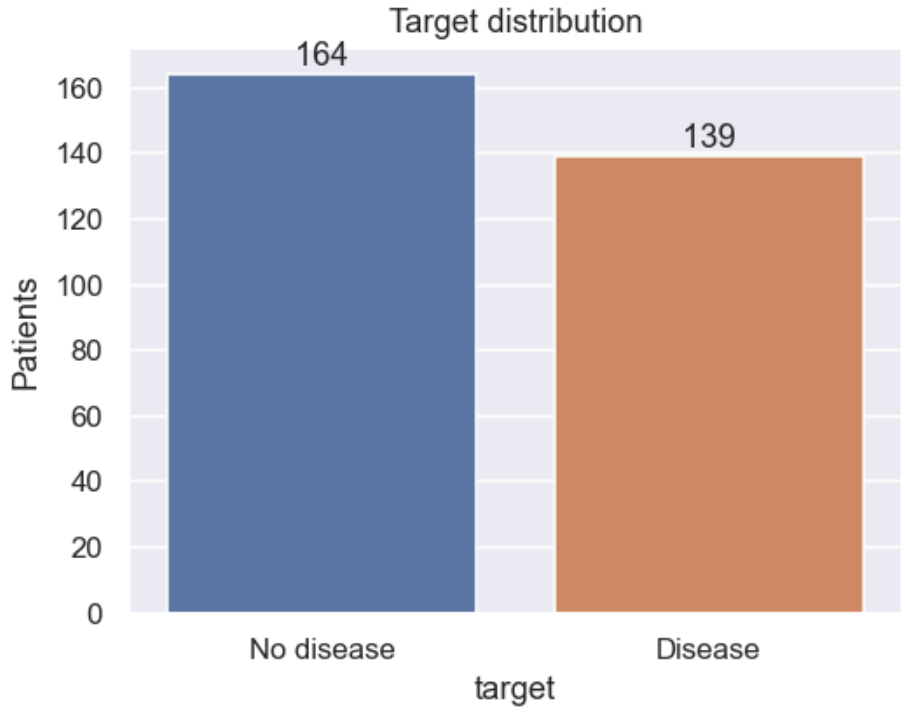
2. Dataset & EDA

Dataset: UCI Cleveland Heart Disease dataset — 303 patients, 13 features, binary target (collapsed from the original 0-4 severity score). Fetched via `heart_risk.ingestion.ClevelandHeartDiseaseIngestor`, a small dataclass rather than a loose script, so the source URL / output path are grouped as configuration. Not committed to git — re-fetched on demand in ~2 seconds.

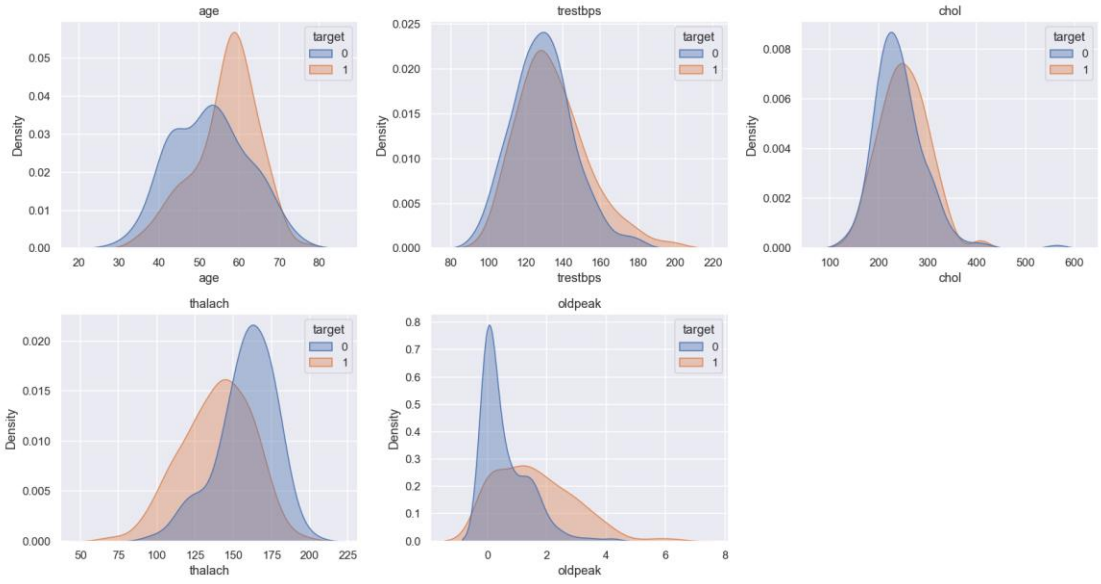
Missing data: 6 of 303 rows (2%) have gaps, confined to `ca` and `thal`. Handled with `KNNImputer` for continuous columns (see Part II) rather than dropping rows.

Key EDA findings (full notebook: `notebooks/exploratory_analysis.ipynb`):

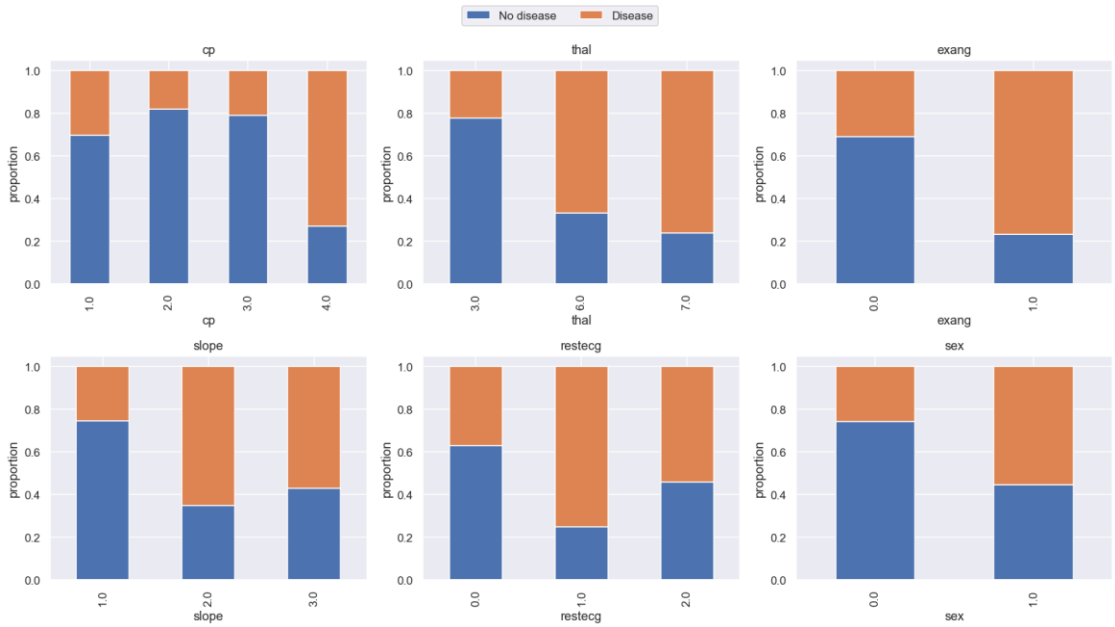
- **Class balance:** 46% positive / 54% negative — mild imbalance, handled with stratified splitting throughout rather than resampling.
- **`thalach`** (max heart rate) shifts lower and **`oldpeak`** (ST depression) shifts higher in the disease group — both visible clearly in the KDE overlays.
- **`thal`** and **`cp`** show the strongest categorical separation — certain codes are almost entirely one class.
- **No problematic collinearity** between predictors — nothing dropped for redundancy.



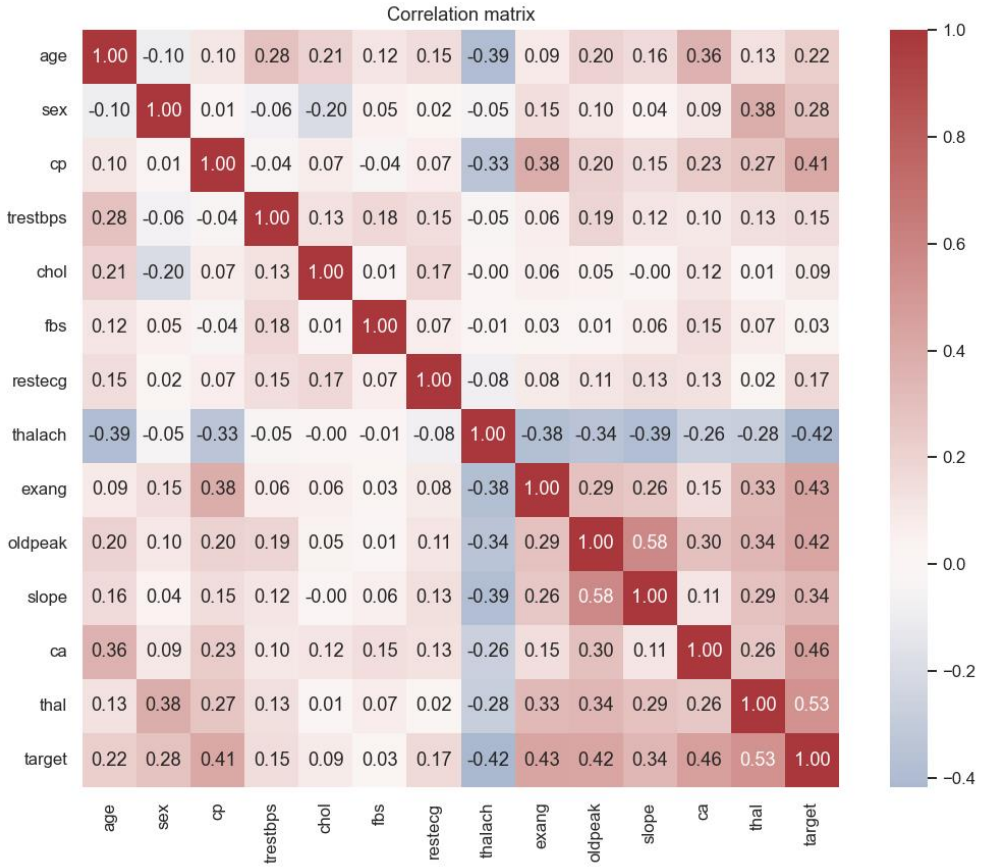
Target distribution



Continuous features by target (KDE)



Categorical feature proportions by target



Correlation matrix

Part II — Modelling

3. Feature Engineering

Implemented in `heart_risk/transform.py` as a two-stage pipeline: an `AgeGroupBucketer` transformer (adds a derived `age_group` column, bucketing age into young/middle/senior), feeding into a `ColumnTransformer`:

Feature group	Columns	Transform	Why
Continuous	age, trestbps, chol, thalach, oldpeak	`KNNImputer(k=5)` -> `StandardScaler`	Only 6 missing cells total — KNN imputation is a near-zero-cost upgrade over median at this data size
Nominal categorical	cp, restecg, slope, thal, ca	most-frequent impute -> `OrdinalEncoder`	Both models are tree-based; integer-coded categories don't impose a false distance metric for trees the way they would for a linear model, and this keeps the feature space compact on a 303-row dataset
Binary	sex, fbs, exang	most-frequent impute	Already 0/1
Engineered	age_group (derived)	pass-through	Clinically motivated bucket (young <45 / middle 45-59 / senior 60+)

This produces **14 output features** (13 raw + 1 engineered), versus what a one-hot approach would produce (20+) — deliberately compact given the small sample size.

4. Model Training & Results

Models trained (`heart_risk/modeling/train.py`): a single **Decision Tree** and a **Gradient Boosting** ensemble, each tuned via `RandomizedSearchCV` (10-fold **stratified** CV, 20 sampled candidate configurations, scored on **F1** rather than ROC-AUC — chosen because F1 balances precision/recall directly, which matters for a screening tool where both false negatives and false positives carry cost).

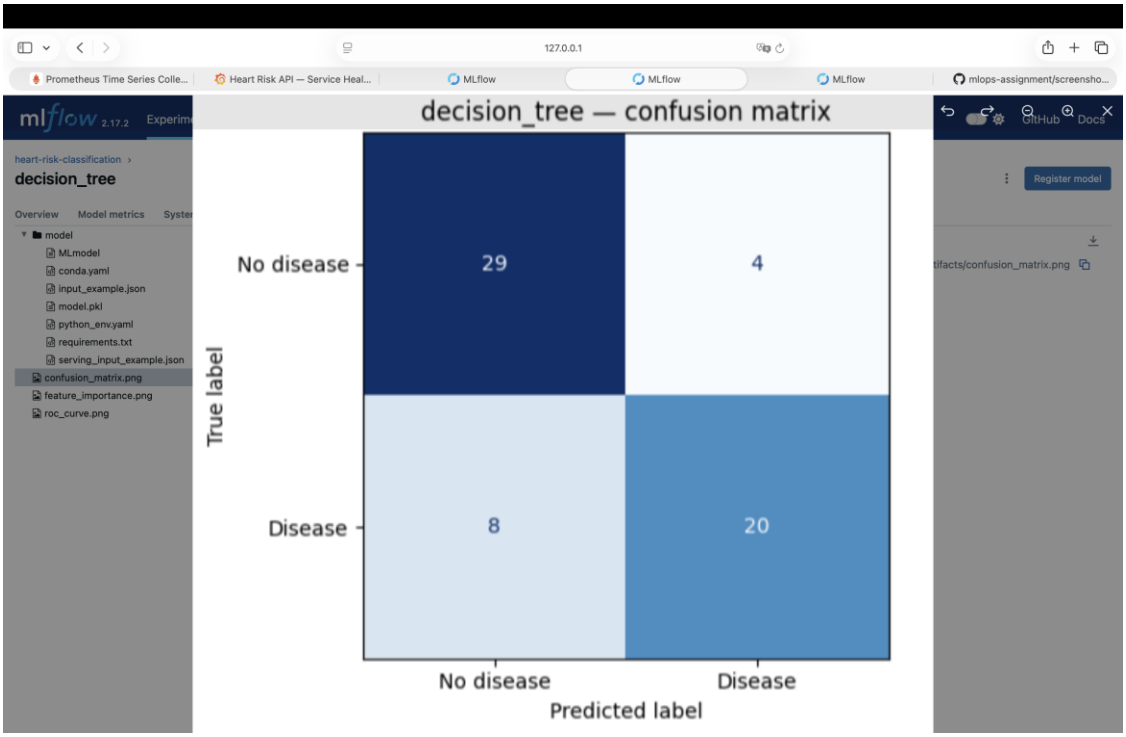
- Decision Tree search space: `max_depth`, `min_samples_split`, `min_samples_leaf`, `criterion`
- Gradient Boosting search space: `n_estimators`, `learning_rate`, `max_depth`, `subsample`

Results

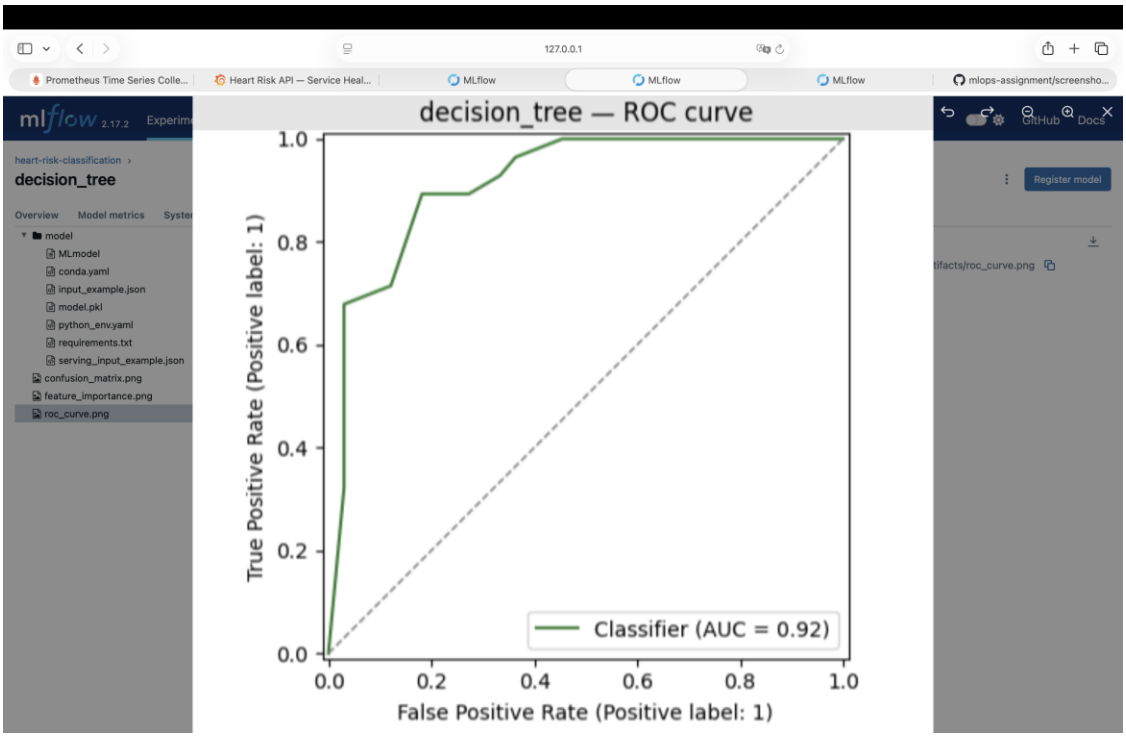
Model	CV F1 (10-fold)	Holdout Accuracy	Holdout Precision	Holdout Recall	Holdout F1	Holdout ROC-AUC
Gradient Boosting (champion)	0.799	0.820	0.870	0.714	**0.784**	**0.939**
Decision Tree	0.736	0.803	0.833	0.714	0.769	0.917

Champion hyperparameters: `n_estimators=50`, `learning_rate=0.05`, `max_depth=2`, `subsample=0.85`.

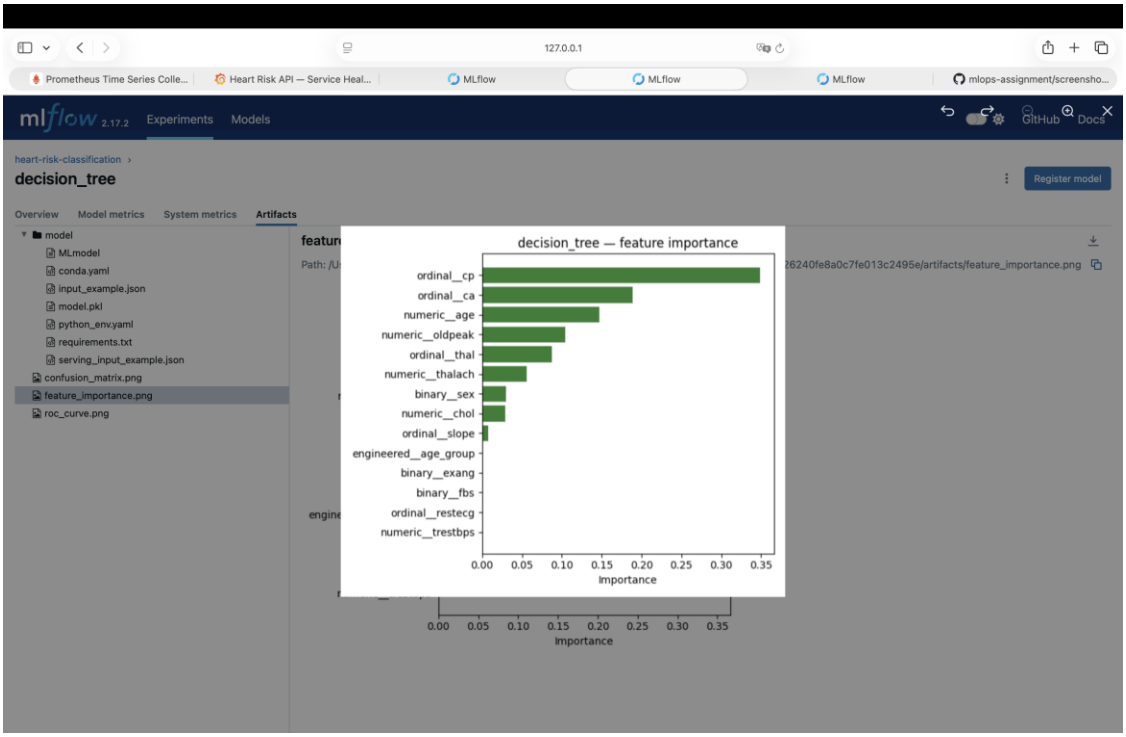
Why Gradient Boosting won: it outperforms the single Decision Tree on every holdout metric except recall (tied). This tracks with expectations — boosting an ensemble of shallow trees (max_depth=2 here) typically generalizes better than one deeper tree on tabular data this size, without the variance a single unconstrained tree is prone to. The feature-importance plots confirm `thal`, `cp`, and `oldpeak` (labelled `ordinal__cp`, `ordinal__ca`, `numeric__oldpeak`, `ordinal__thal` after the preprocessing pipeline) as the top drivers for both models, consistent with the EDA.



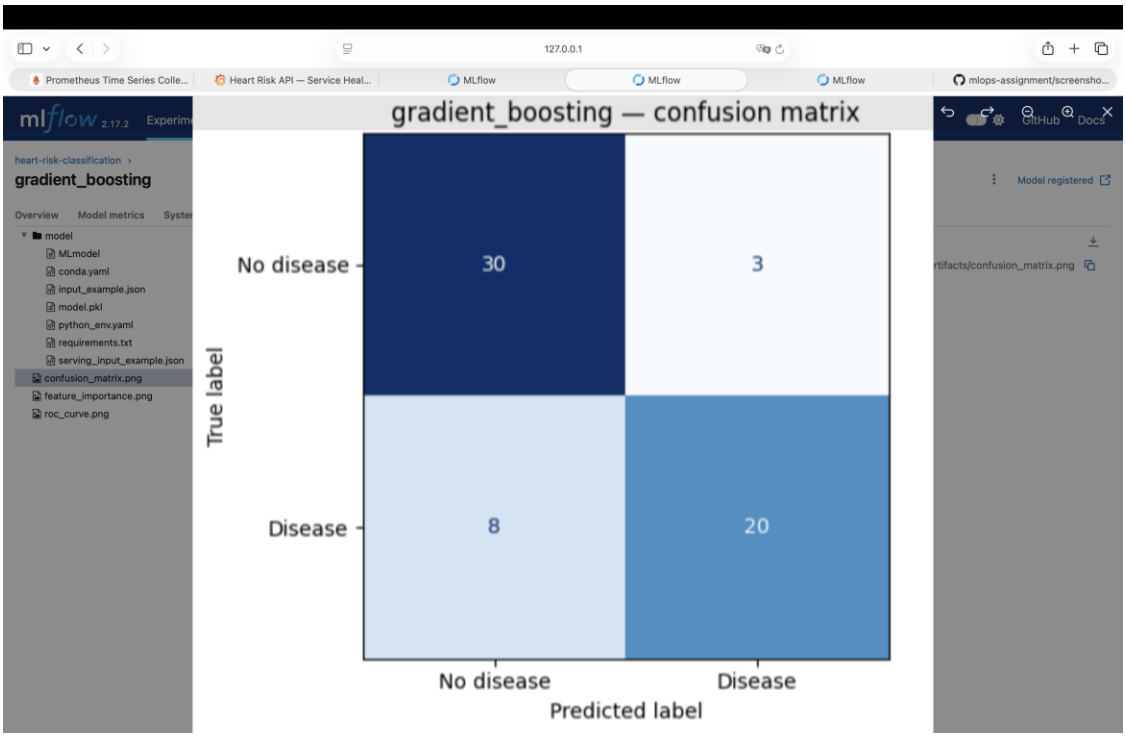
Decision Tree confusion matrix



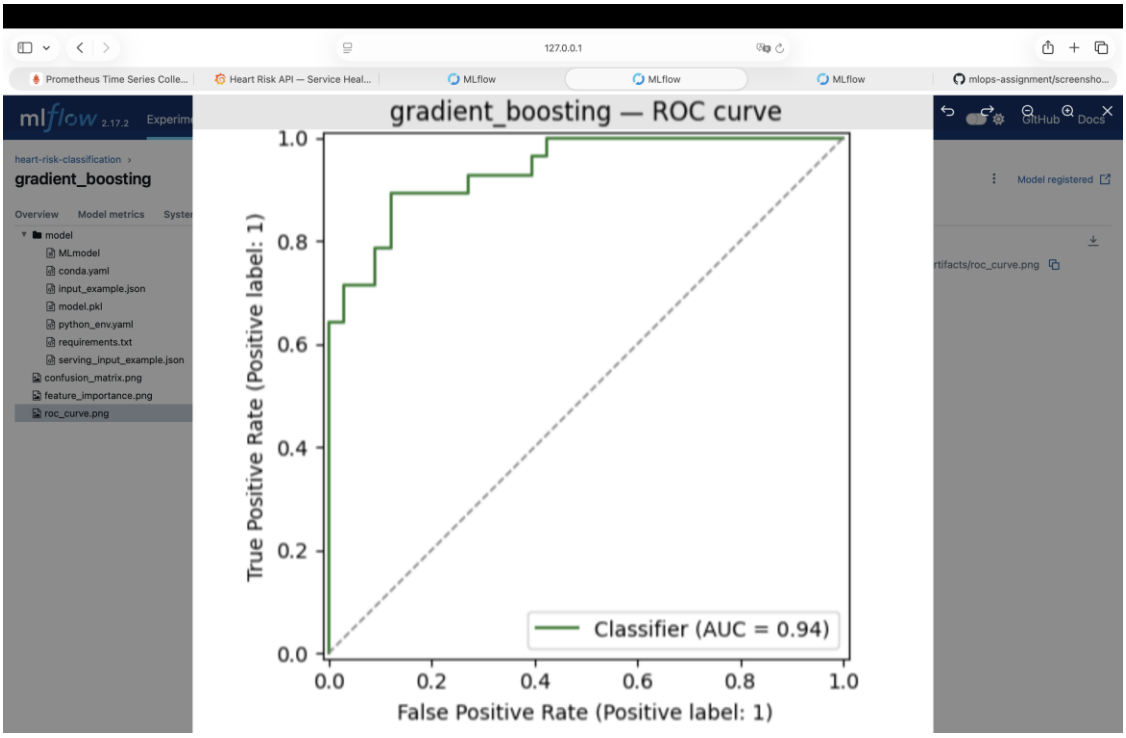
Decision Tree ROC curve



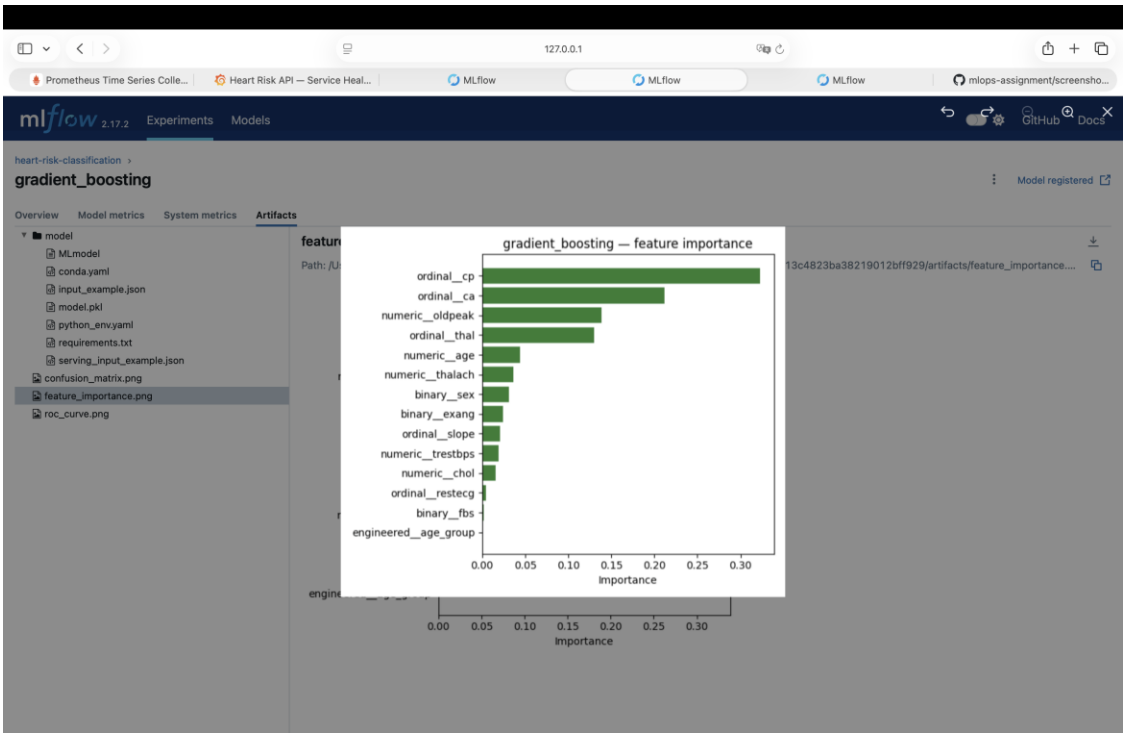
Decision Tree feature importance



Gradient Boosting confusion matrix



Gradient Boosting ROC curve



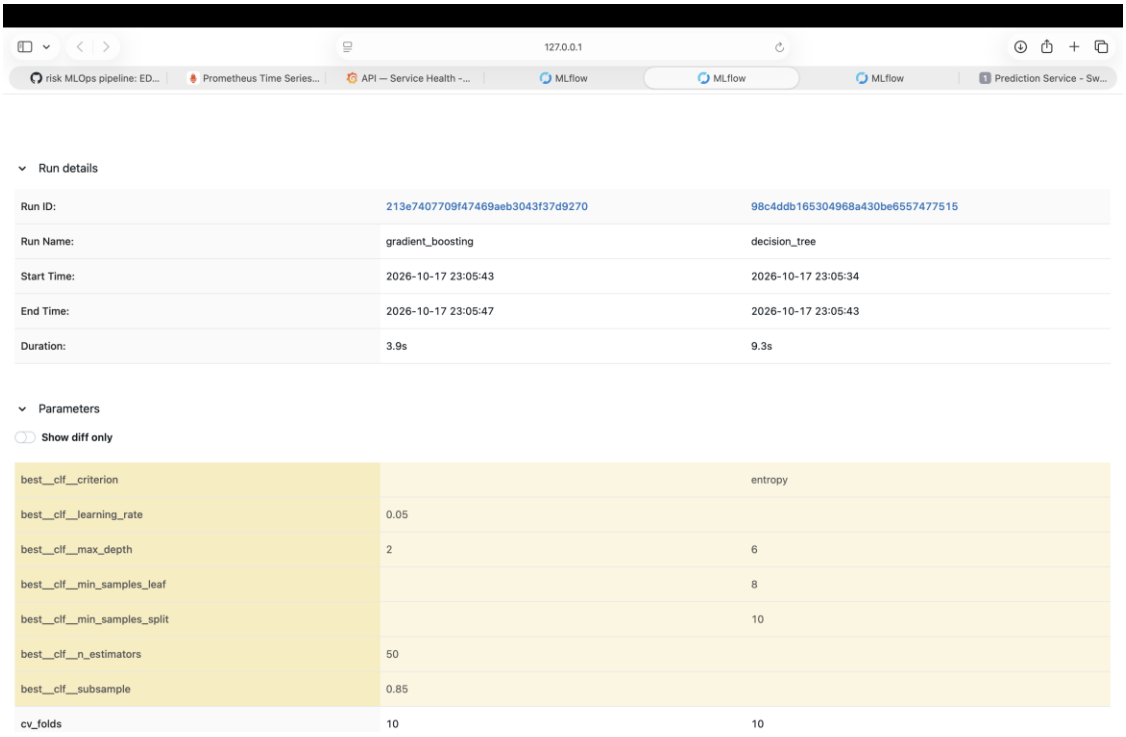
Gradient Boosting feature importance

5. Experiment Tracking (MLflow)

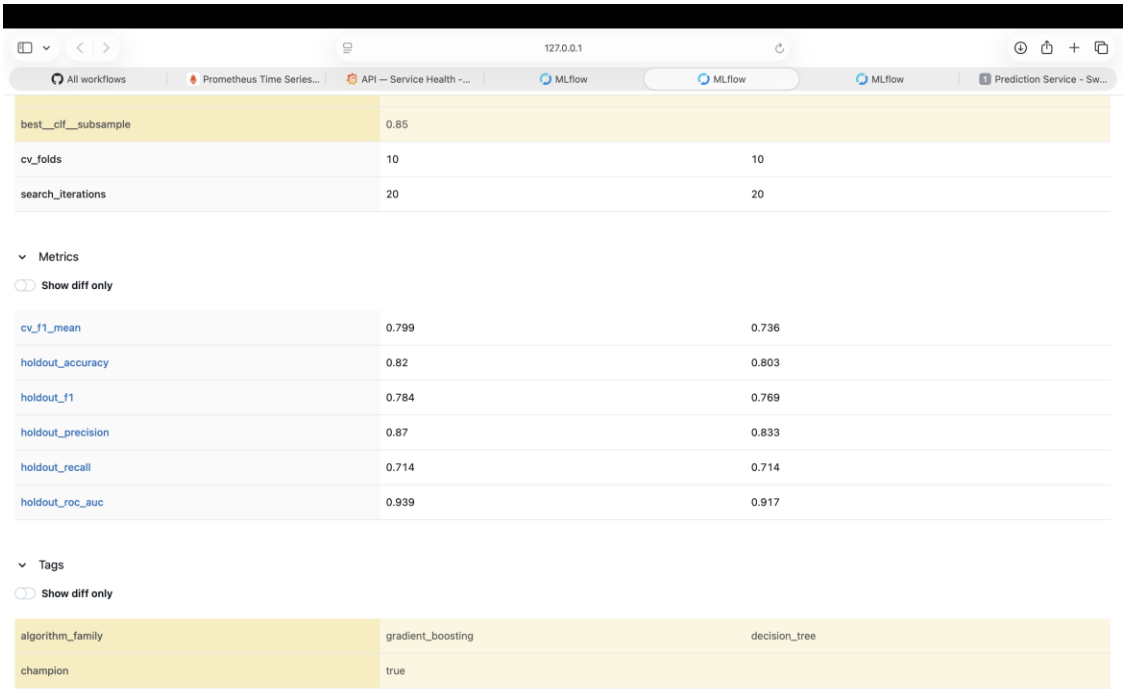
Tracking store: **`sqlite:///mlflow.db`** (not the default local file store) — specifically so the champion model could be registered in the **MLflow Model Registry**, which requires a database-backed store. Every `RandomizedSearchCV` run is logged under the `heart-risk-classification` experiment with:

- **Params:** winning hyperparameters, CV fold count, search iteration count
- **Metrics:** CV F1 mean, and holdout accuracy/precision/recall/F1/ROC-AUC
- **Artifacts:** confusion matrix, ROC curve, **and a feature-importance plot** (both models expose `feature_importances_` natively) — an artifact a Logistic-Regression-based pipeline wouldn't give for free
- **Registry:** the champion is registered as `heart-risk-classifier`, giving it a version number independent of the underlying MLflow run

Side-by-side run comparison from the MLflow UI (`gradient_boosting` vs. `decision_tree`, matching the results table above exactly, including the `champion=true` tag set on the winner):



MLflow run comparison - parameters and tags



MLflow run comparison - metrics

6. Model Packaging & Reproducibility

- The champion pipeline is saved via **`mlflow.sklearn.save_model`** — an MLflow-format model *directory* (`artifacts/model/`, containing `MLmodel`, `model.pkl`, `conda.yaml`, `requirements.txt`) rather than a bare `joblib.dump` pickle. The serving layer loads it with `mlflow.pyfunc.load_model`, so swapping in a different MLflow-flavored model later wouldn't require touching the loading code.
- `artifacts/metadata.json` records the champion's name, registry version, hyperparameters, holdout metrics, and a summary of every candidate trained.
- Two requirements files: `requirements.txt` (full dev/training environment) and `deploy/docker/requirements.txt` (slim serving-only set using **`mlflow-skinny`** instead of full `mlflow`, to avoid pulling the tracking server/UI stack into the serving image). All packages version-pinned.

Part III — Engineering & Operations

7. Testing Strategy

Unit tests (`tests/`, 41 cases, Python's built-in **`unittest`** framework — `TestCase` classes with `setUp`/`tearDown`, chosen over pytest as the assignment's other explicitly listed option): cover ingestion parsing, dataset loading/splitting, the feature pipeline (no-NaN output,

ordinal-not-onehot width check, independent refits), the `AgeGroupBucketer`, metric computation + all three plot functions, the training candidates' structural correctness, `RiskPredictor` inference, and the FastAPI service's endpoints. No test touches the network or depends on `artifacts/model/` already existing — a synthetic fixture dataset (`tests/fixtures.py`) stands in for the real data.

8. CI/CD Pipeline

`.github/workflows/pipeline.yml` — three Makefile-driven jobs, each gated on the previous succeeding:

1. **quality** — `make lint` (flake8), `make test` (unittest)
2. **train** — `make ingest`, `make train`; asserts `artifacts/metadata.json` was produced; uploads the model + MLflow SQLite DB as an artifact
3. **containerize** — downloads the trained model, `make docker-build`, `make docker-run`, polls `/health`, sends a real `/predict` request and asserts the response shape, then **lints the Helm chart** (`make helm-lint`)

Any failure — lint error, failing test, training crash, a container that never becomes healthy, a bad `/predict` response, or an invalid Helm chart — fails the pipeline with a non-zero exit and a readable log.

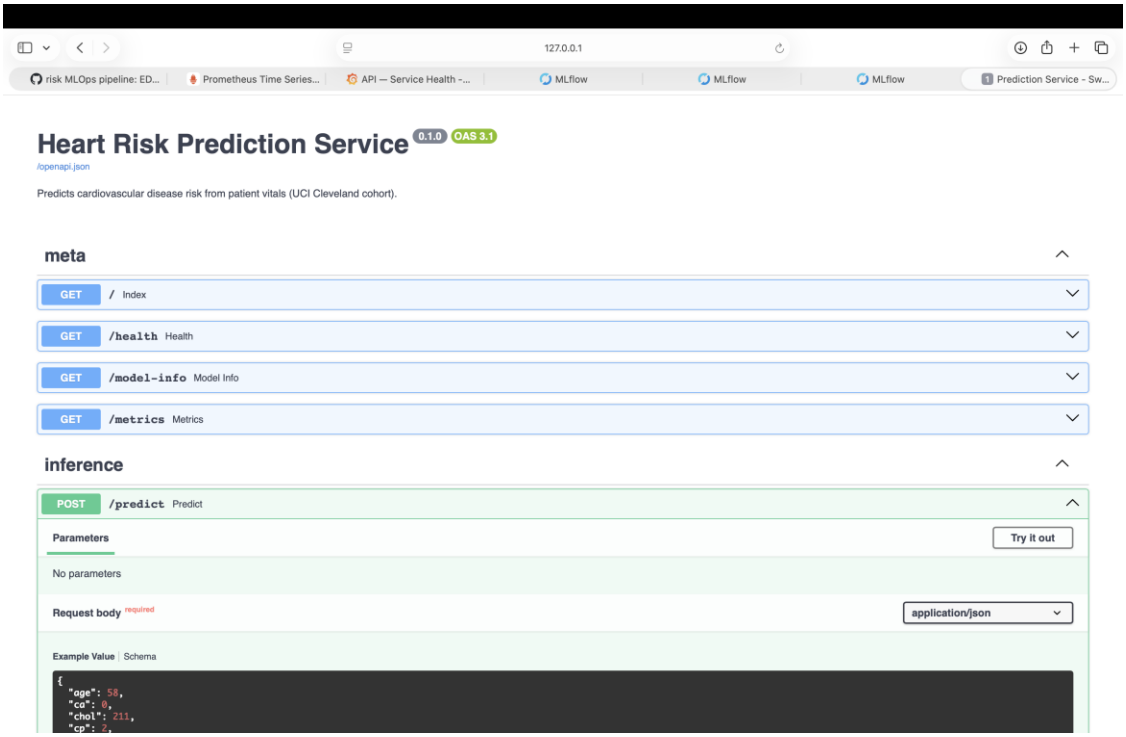
9. Containerization

`deploy/docker/Dockerfile` — multi-stage build: - **Builder stage:** Python 3.12-slim + `build-essential`, installs `deploy/docker/requirements.txt` into an isolated venv - **Runtime stage:** Python 3.12-slim + the venv copied over, app code, and the MLflow-format model directory — no compiler toolchain in the final image. Runs as a non-root user, `HEALTHCHECK` against `/health`. - **Process model:** `gunicorn` managing `uvicorn.workers.UvicornWorker` processes (2 workers) rather than invoking uvicorn directly — closer to how this would actually run in production, and gives explicit worker-count control.

Verified locally:

```
make docker-build
make docker-run
curl http://localhost:8080/health
# {"status":"ok","model_ready":true,"model_name":"heart-risk-classifier","mod
el_version":"1"}
```

The FastAPI interactive docs (Swagger UI), auto-generated from the Pydantic contracts in `heart_risk/service/contracts.py`:



FastAPI Swagger UI - Heart Risk Prediction Service

10. Production Deployment (Kubernetes via Helm)

Packaged as a **Helm chart** (`deploy/helm/heart-risk-api/`) rather than raw manifests — the assignment explicitly allows either, and a chart makes the "bundle the API + its monitoring stack as one installable unit" design (Section 11) much cleaner than separate `kubectl apply` files.

- **Deployment:** 2 replicas, resource requests/limits, readiness probe (`/health`, 5s delay) and liveness probe (`/health`, 15s delay)
- **Service:** `ClusterIP` (not directly exposed)
- **Ingress:** nginx ingress class, host `heart-risk.local` — the assignment's *other* explicit option for exposing the API (vs. a LoadBalancer Service)

```
minikube addons enable ingress
helm install heart-risk-api deploy/helm/heart-risk-api \
  --namespace heart-risk --create-namespace
```

Verified end-to-end through the Ingress (not just the pod's own loopback):

```
$ kubectl -n heart-risk get pods,svc,ingress
heart-risk-api-...   1/1   Running   (x2 replicas)
heart-risk-api-grafana-...      1/1   Running
heart-risk-api-prometheus-...   1/1   Running
ingress.networking.k8s.io/heart-risk-api   nginx   heart-risk.local   192.168.
49.2
```

```
$ curl -H "Host: heart-risk.local" http://<ingress-ip>/health
{"status":"ok","model_ready":true,...}

$ curl -H "Host: heart-risk.local" -X POST .../predict -d '{...}'
{"prediction":0,"risk_label":"low_risk","confidence":0.862,...}
```

```
NAME                                          READY   STATUS    RESTARTS   AGE
heart-risk-api-c99d88478-5vzxf                1/1     Running   0          56m
heart-risk-api-c99d88478-rvk2t                1/1     Running   0          56m
heart-risk-api-grafana-94c7c44bc-hbs6m        1/1     Running   0          24h
heart-risk-api-prometheus-779656b487-kxcx8    1/1     Running   0          24h
```

kubectl get pods -n heart-risk - API (2/2) + Prometheus + Grafana all Running

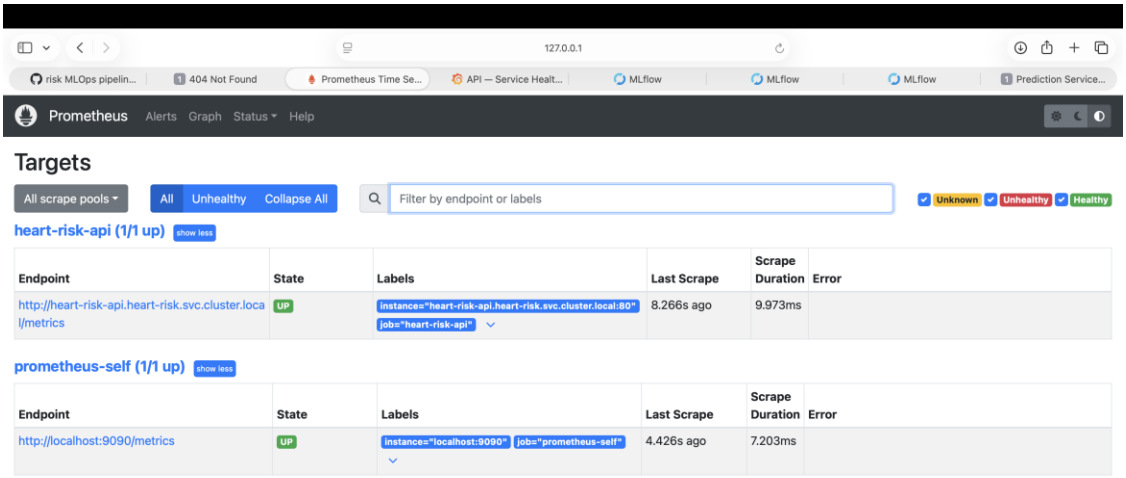
11. Monitoring & Logging

Request logging: every request produces a single **structured JSON log line** (trace id, method, route, status, duration) via a custom `JsonLogFormatter` — closer to what a real log aggregator (Loki, CloudWatch) would expect than a plain f-string.

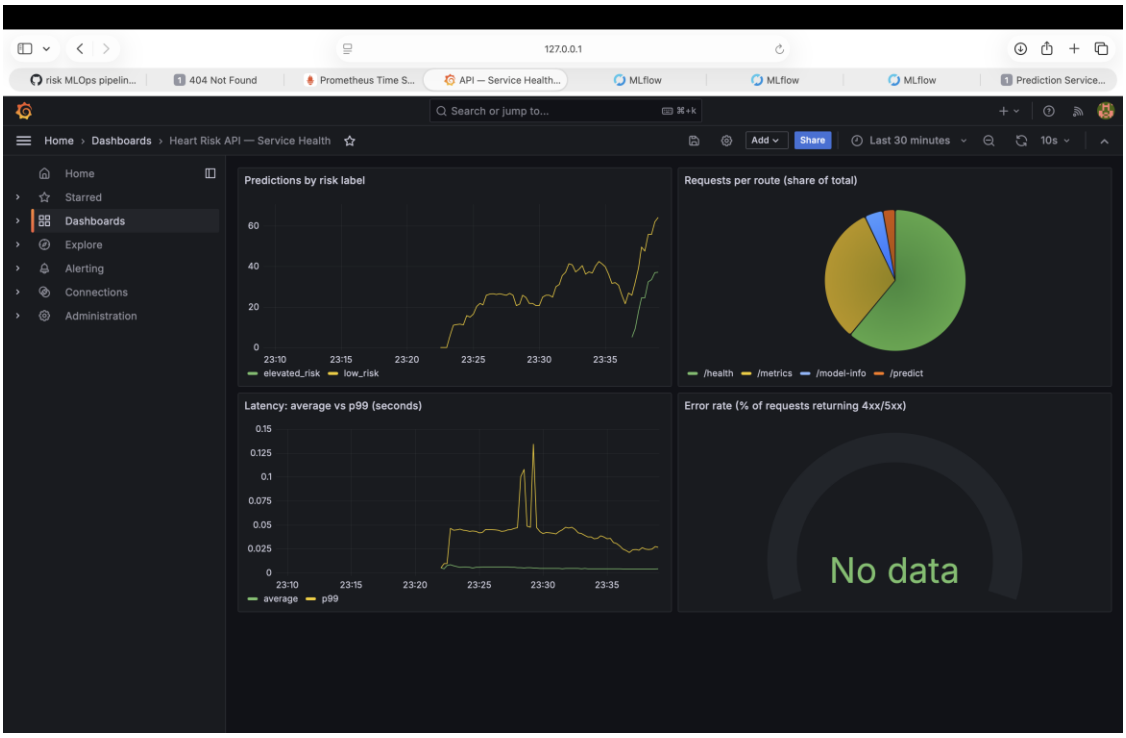
Metrics: namespaced `heart_risk_*` Prometheus metrics (to avoid colliding with other services on a shared Prometheus instance): - `heart_risk_http_requests_total{route,method,status}` - `heart_risk_predictions_total{risk_label}` - `heart_risk_request_duration_seconds{route}`

Deployed stack: Prometheus and Grafana are templated as part of the **same Helm chart** as the API (`deploy/helm/heart-risk-api/templates/monitoring/`), so `helm install` brings up the entire observability stack alongside the service — no separate manifests to apply. Grafana auto-provisions: - A Prometheus datasource - A **"Heart Risk API — Service Health"** dashboard with four panels deliberately varied in type: a predictions-by-label **time series**, a requests-per-route **pie chart**, an average-vs-p99 latency **time series**, and an error-rate **gauge** — not just four line charts of different metrics.

Verified live against real traffic generated through the Ingress: Prometheus reports the `heart-risk-api` target as **UP**, and Grafana's API confirms both the datasource and dashboard are provisioned.



Prometheus targets - heart-risk-api UP



Grafana - Heart Risk API Service Health dashboard with live data

Appendix

A. Setup Instructions

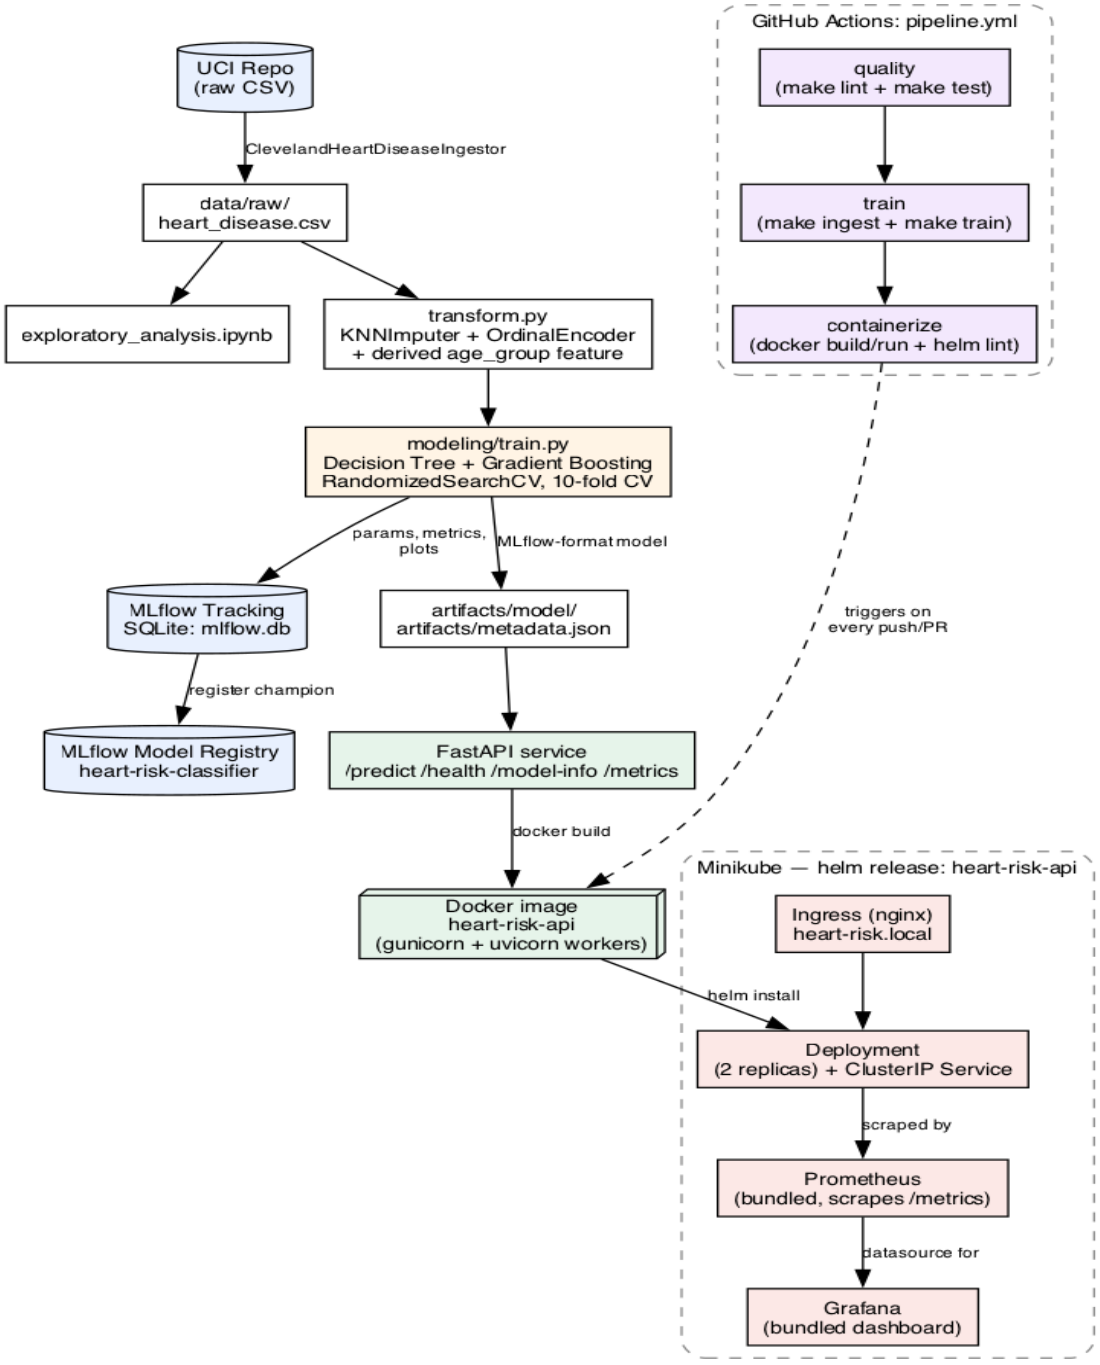
```
git clone https://github.com/arya2305/mlops-workflow.git
cd mlops-workflow

python3.12 -m venv .venv
source .venv/bin/activate
make install

make ingest     # fetches data/raw/heart_disease.csv from UCI
make train      # trains + tracks in MLflow, saves artifacts/model/
make test       # 41 unittest cases, ~4s
make serve      # API on http://localhost:8080
```

Full step-by-step instructions for Docker, Kubernetes/Helm, and monitoring are in the project `README.md`.

B. Architecture Diagram



Architecture diagram

The authoritative, version-controlled source for this diagram is a Mermaid flowchart in `README.md` (renders natively on GitHub); the PNG above is a matching static rendering for this Word/PDF report.

C. Repository Link

https://github.com/arya2305/mlops-workflow